

Waddington: a hybrid ML + LLM agent for sequential CRISPR gene selection

Results tables and figures

All numbers in this document are computed directly from the frozen benchmark outputs (the JSON files under `workspace/results/sequential/`) by `waddington_select.analysis`; no value is transcribed by hand. Figures are produced by `python -m waddington_select.analysis figures`.

Table 1: The nine benchmark screens. The phenotype scores and hit sets are taken *unchanged* from the BioDiscoveryAgent (BDA) release (Roohani et al., ICLR 2025); we verified for all nine that our hit labels are exactly `topmovers` $\cap$ (pool $\setminus$ CEGv2). Task descriptions are BDA’s own task prompts, not our paraphrases. **We remove the 684 core-essential genes of CEGv2** (Hart et al.) from the candidate pool — BDA’s own non-essential (“N/E”) convention, since essential genes are detected as hits under almost any screen and reward nothing. Columns show the pool *before* $\rightarrow$ *after* that filter; all results use the filtered pool (the two K562 screens contain no CEGv2 gene, so they are unchanged).

Perturbation. KO = CRISPR knockout (loss of function); CRISPRi = interference/knockdown (loss of function); CRISPRa = activation (gain of function). Note that our DepMap features are *all knockout-derived*, so they transfer least well to the gain-of-function screen (Steinhart) — which is exactly where the LLM contributes most (Table 5). *The Schmidt IL-2 and IFN- γ data are the *CRISPRi* (loss-of-function) arm. Schmidt et al. ran both CRISPRa and CRISPRi genome-wide; BDA sources these two screens from GeneDisco, whose benchmark explicitly uses the *interference* screens (Mehrjou et al., 2021, §4.2.2–4.2.3), although BDA’s own task prompt does not state the arm. This matters here: it makes all eight non-Steinhart screens loss-of-function, leaving Steinhart the lone gain-of-function screen (Table 6). ‡The two Perturb-seq screens define a hit as a strong *transcriptome-wide* response (top 10%), not a specific phenotype at the top 5% used for the other seven — which is why their hit rates are higher.

Steinhart is not a published study: it is unpublished data generated by a BDA co-author (Z. Steinhart, Marson lab), which BDA included specifically as a leakage control — “one dataset that is unpublished, ensuring it is not part of the language model’s training data”. That control was valid for *their* models but **not for ours**: BDA released the screen in their public repository in 2024, before our LLM’s training cutoff. See Limitations.

Dataset	Task (BDA prompt)	Perturb.	all $\rightarrow$ non-essential		Hit rate	Batch
			Genes	Hits		
IFN- γ	IFN- γ production*	CRISPRi	18,418 $\rightarrow$ 17,785	920 $\rightarrow$ 802	4.5%	128
IL-2	IL-2 production*	CRISPRi	18,939 $\rightarrow$ 18,273	654 $\rightarrow$ 481	2.6%	128
Sanchez21	Tau change, <i>either</i> direction	KO	18,469 $\rightarrow$ 17,807	924 $\rightarrow$ 859	4.8%	128
Sanchez21 (dn)	Tau <i>decrease</i>	KO	18,469 $\rightarrow$ 17,807	924 $\rightarrow$ 871	4.9%	128
Carnevale22	T-cell efficacy (adenosine)	KO	18,861 $\rightarrow$ 18,224	943 $\rightarrow$ 868	4.8%	128
Scharenberg22	Lysosomal choline recycling	KO	1,061 $\rightarrow$ 1,029	53 $\rightarrow$ 49	4.8%	32
Steinhart	Resisting T-cell exhaustion	CRISPRa	18,797 $\rightarrow$ 18,144	149 $\rightarrow$ 145	0.8%	128
K562-Essential	Strong transcriptomic response‡	CRISPRi	623 $\rightarrow$ 623	63 $\rightarrow$ 63	10.1%	32
K562-GWPS	Strong transcriptomic response‡	CRISPRi	9,193 $\rightarrow$ 9,193	924 $\rightarrow$ 924	10.1%	128

Relationship to BioDiscoveryAgent

We use BDA’s *data*, not their method: the ground-truth phenotype scores and `topmovers` hit sets are theirs, and we verified that our hit labels are exactly `topmovers` $\cap$ (pool $\setminus$ CEGv2) for every screen. **We did not run BioDiscoveryAgent, and the numbers in this paper are not a reproduction of theirs and are not directly comparable to them.** Two differences matter:

- **We add a cross-experiment ML prior that BDA does not have.** Our LOO-LightGBM is trained on the *other* screens’ hit labels using PPI / KEGG / pLI / DepMap features. Restricted to the six screens we share with BDA, that model alone averages 0.187 (Table 2) — already well above BDA’s best agent (0.127) and the best GeneDisco baseline (0.091),

with no LLM at all. Our gains over BDA’s published numbers should therefore be read as “a cross-experiment prior helps a lot”, not as “our agent is a better agent”.

- **Different LLM and different baseline implementations** (we use Claude Haiku 4.5; BDA’s headline model is Claude 3.5 Sonnet; our Coreset is our own implementation and scores differently from theirs, e.g. 0.100 vs. 0.066 on IFN- γ).

As a partial check that the conventions align, our *random* baseline agrees on some screens (IL-2 0.031 vs. 0.031; Steinhart 0.020 vs. 0.021) but — because we re-derive the hit sets from the same top-5% $\cap$ non-CEGv2 convention rather than copy them — differs by up to 0.034 on others (Scharenberg22 0.102 vs. 0.136). Table 2 is therefore an *approximate* reference, not a controlled head-to-head; it holds up only because the gaps dwarf that imprecision.

Table 2: External reference on the six screens BDA and we share (Schmidt1/2 = IFN- γ /IL-2, CART = Steinhart, plus Scharenberg22, Carnevale22, Sanchez21). Mean hit@R5 over the non-essential (N/E) evaluation, 5 rounds, same batch sizes throughout. **Prior-work rows are transcribed from Roohani et al. (2025), Table 1 — we did not re-run them.** “GeneDisco best-per-screen” is an oracle taking the strongest of their eight baselines on each screen. Our random differs from theirs by up to 0.034 (see above), so treat gaps below ~ 0.03 as noise.

Method	Source	mean hit@R5
Random	BDA (transcribed)	0.049
Random	ours	0.041
GeneDisco — best single baseline (Coreset)	BDA (transcribed)	0.083
GeneDisco — best-per-screen (oracle over 8)	BDA (transcribed)	0.091
BioDiscoveryAgent — Claude 3.5 Sonnet	BDA (transcribed)	0.127
LOO-LightGBM — no LLM (ours)	ours	0.187
Waddington C (ours)	ours	0.217

The *ordering* is the point, not the exact values. Our cross-experiment prior alone — a static LightGBM with no LLM — exceeds BioDiscoveryAgent’s strongest agent by +0.060 ($\sim 7\times$ the random discrepancy), and BDA’s best in turn roughly doubles the strongest GeneDisco baseline. Per the caveats above, our lead is an *information* advantage: the LOO model sees the other screens’ hit labels, which neither GeneDisco nor BDA has. The honest reading is a hierarchy of *priors* — diversity / uncertainty heuristics (GeneDisco) < a strong LLM’s parametric biology (BDA) < a cross-experiment supervised prior (ours) — not “our agent beats their agent”.

Table 3: Hit ratio at round 5 across all nine benchmark datasets. All methods are reported over 5 seeds. **Bold** = best per row; underline = second best. Waddington is our proposed method (C-arm), using DeLM-verified cross-experiment memory.

[†]**B is not a pure-LLM baseline.** When the LLM names too few genes that exist in the pool, the batch is padded from the static LOO-LightGBM ranking. Averaged over the five rounds that padding is 19% (IFN- γ), 19% (IL-2), 46% (Sanchez21), 32% (Sanchez21-dn), 7% (Carnevale22), 68% (Scharenberg22), 43% (Steinhart), **86% (K562-Essential)** and 0% (K562-GWPS). B should be read as “LLM, padded with a static ML prior”, and its two strongest entries (K562-Essential 0.559, Scharenberg22 0.478) are largely that prior rather than the LLM.

Dataset	Random	LOO-LightGBM	OnlineAdapt.	A: Coreset	B: LLM [†]	C: Waddington
IFN- γ	0.029	0.168	<u>0.183</u>	0.100	0.160	0.190
IL-2	0.031	0.306	<u>0.314</u>	0.139	0.257	0.347
Sanchez21	0.037	0.077	<u>0.087</u>	0.031	0.063	0.087
Sanchez21 (dn)	0.029	0.091	0.101	0.078	0.071	<u>0.099</u>
Carnevale22	0.024	0.048	0.047	0.054	0.059	<u>0.056</u>
Scharenberg22	0.102	0.449	0.449	0.286	0.478	<u>0.465</u>
Steinhart	0.021	0.076	0.090	0.090	<u>0.152</u>	0.159
K562-Essential	0.254	0.492	0.476	0.270	0.559	<u>0.556</u>
K562-GWPS	0.065	0.247	<u>0.273</u>	0.156	0.225	0.347
Average	0.066	0.217	0.224	0.134	<u>0.225</u>	0.256

Waddington (C) achieves the best average hit ratio (0.256), outperforming the LLM baseline B by 13.7% relative (+0.031) and the algorithmic baseline A by 91.6% (+0.122). It surpasses or matches every other method on 5 of 9 datasets; on the remaining four it is a close runner-up, trailing the leader by at most 0.013. Notably, Waddington strongly exceeds OnlineAdaptive (the PerTurboAgent-style ML-only method) on K562-GWPS (+0.075) and IFN- γ (+0.007). Note that on the two screens where B edges ahead of C (Scharenberg22, K562-Essential) the LLM is barely doing the work: 68% and 86% of B’s batch there is static-ML padding (Table 3[†]), so those entries do not support a claim about LLM reasoning.




Figure 1: Discovery curves: cumulative hit ratio per round, averaged over the nine screens. The hybrid C-arm separates from every baseline from round 2 onwards. Per-screen values are given in Figure 3; no error band is drawn because the spread across screens is dominated by how hard each screen is (K562-Essential 0.56 vs. Carnevale22 0.06), which would encode screen difficulty rather than uncertainty about a method.




Figure 2: Left: average hit ratio @ R5 per method (Table 3). Right: paired ablation deltas (Table 5). Colour states a fact about the *variant* (it scored lower than the full arm), not a verdict on the component: a red “– online ML” means *removing* online ML costs 0.012, i.e. it is the most valuable part.




Figure 3: Per-screen hit ratio at round 5 for every method.

Table 4: Stepwise improvement in average hit ratio at round 5 (9 datasets, 5 seeds). Each row introduces one additional component over the previous configuration.

Configuration	Key addition	avg hit@R5	Δ
Random	—	0.066	—
Coreset (A)	Diversity-maximising selection	0.134	+0.068
LOO-LightGBM (static)	Cross-expt. ML prior	0.217	+0.083
OnlineAdaptive	In-expt. ML retraining	0.224	+0.007
LLM Reasoning (B)	LLM biological prior (no ML)	0.225	(<i>parallel branch</i>)
C – ML (static LOO + LLM)	Combine LOO prior with LLM	0.247	+0.022 vs. B
Waddington (C, ours)	Add online ML retraining	0.256	+0.009

Each component provides an additive gain: the cross-experiment LOO prior (+0.083) contributes most, followed by the LLM biological prior (+0.022 when combined with the LOO prior), and online ML retraining (+0.009). Cross-experiment memory contributes exactly 0.000 in this leave-one-out benchmark (paired ablation, Table 5), as the LLM’s parametric knowledge already subsumes the explicit memory entries.

Table 5: Ablation study, reported as **paired** deltas $\Delta = \text{variant} - \text{C}$, where C is the full arm *re-run inside the same experiment* (same seeds, same LLM sampling). This matters: the arm is stochastic, so the C baseline re-runs to 0.259–0.262 across the ablation experiments rather than exactly 0.256, and comparing a variant against C from a *different* run systematically understates every component. **C – Mem**: cross-experiment memory removed from the LLM prompt. **C – LLM**: LLM removed; online ML only. **C – ML**: ML frozen at the LOO prior; no in-experiment retraining. **Shuffled**: gene names shown to the LLM replaced with anonymous identifiers (GENE.00001, ...). **Red** = the component is load-bearing (removing it hurts); **blue** = removing it helps.

Dataset	Δ C – Mem	Δ C – LLM	Δ C – ML	Δ Shuffled
IFN- γ	+0.003	+0.000	−0.013	−0.015
IL-2	−0.013	−0.002	+0.002	−0.005
Sanchez21	+0.004	+0.010	−0.010	−0.008
Sanchez21 (dn)	−0.005	−0.003	−0.010	−0.004
Carnevale22	−0.002	−0.003	+0.006	+0.001
Scharenberg22	+0.012	+0.114	−0.016	+0.078
Steinhart	−0.008	−0.059	−0.004	−0.084
K562-Essential	+0.016	−0.127	−0.025	−0.051
K562-GWPS	−0.004	−0.004	−0.038	−0.004
Average (paired)	+0.000	−0.008	−0.012	−0.010

Memory ($\Delta = +0.000$). Removing the cross-experiment memory prompt changes nothing: the paired average is exactly zero. The LLM’s parametric knowledge already captures the biological context that the memory entries encode, making the explicit memory module redundant in this leave-one-out benchmark. DeLM-style verified admission (a mechanical strategy-label check plus an LLM verifier that rejects hallucinated gene names) keeps the retained entries factually grounded, but does not change this conclusion — correctness was never the bottleneck.

Online ML retraining ($\Delta = -0.012$). Freezing the ML model at the LOO prior produces the most consistent degradation of any ablation: it hurts on 7 of 9 screens, with the largest losses on K562-GWPS (−0.038) and K562-Essential (−0.025). Online adaptation is the single most valuable component of the arm.

The “LLM” branch ($\Delta = -0.008$ on average, but bimodal) — and an important caveat. C – LLM removes the entire LLM-branch endorsement term from the fusion. That term is *not* purely the LLM: it also carries the static-ML padding described in Table 3[†]. The two screens with the largest effects are exactly the two with the most padding, so they must be read carefully:

- **K562-Essential** (−0.127): this is *not* evidence for LLM biology. 86% of the branch’s genes there come from the static LOO ranker (the LLM names almost nothing that exists in the 623-gene pool). What the ablation removes is the *static cross-experiment prior* being blended into the online model — a real and valuable signal, but not LLM reasoning. We retract the earlier reading of this number.
- **Scharenberg22** (+0.114 when removed): 68% padding, so “the LLM hurts here” is likewise mostly “the static prior hurts here”, conflicting with an unusually strong online ML signal (LOO AUC = 0.825 on K562 DepMap features).

- **Steinhart (−0.059): this one does survive.** Padding is 43%, so a majority of the branch is genuine LLM naming, and the gene-name ablation below independently confirms it.

Steinhart is also where one would *predict* the LLM to matter most, for a reason visible in Table 1: it is the gain-of-function (CRISPRa) screen, asking which genes make CAR-T cells resist exhaustion *when over-expressed*, whereas every DepMap feature we use is derived from *knockout* screens. A gene that is dispensable when deleted can still be potent when over-expressed — IRF4, one of Steinhart’s strongest hits, is essential in only 6% of DepMap cell lines — so the ML prior has little to say there, and the LLM’s biological knowledge is what fills the gap. Consistently, Steinhart has both the lowest hit rate of the nine screens (0.8%) and the largest gene-name-shuffle penalty (−0.084).

Gene-name semantics ($\Delta = -0.010$; −0.084 on Steinhart) — the clean probe of LLM biology. This ablation is immune to the padding confound in a way C – LLM is not: the branch is still present and still padded, but the LLM can no longer recognise the genes it is naming. Performance collapses on Steinhart (−0.084) and drops on K562-Essential (−0.051), so the LLM’s contribution on pathway-specific tasks is genuinely mediated by *biological gene-name knowledge* rather than by the task description or by experimental feedback. Scharenberg22 again improves (+0.078).

Reasoning vs. recall: can structure replace the gene name?

The shuffle ablation shows the LLM’s contribution runs through gene-name knowledge, but it does not show *why* — because it removes the name and supplies nothing in its place. The shuffled arm must name identifiers out of an 18k space with no information whatsoever; it measures a blindfolded guesser, not a reasoner. The question we actually care about is whether the LLM must *recall* biology from its weights, or whether it can *reason* from structure we hand it.

We therefore ran a four-point decomposition. All four use the identical routing, fusion and update harness; the *only* thing that varies is what the LLM sees:

- **A0** — anonymized IDs, no candidate pool, no features (the shuffle arm above).
- **A0.5** — anonymized IDs + the ML-ranked candidate pool, *no* features (control).
- **A1** — anonymized IDs + the pool + each candidate’s structural feature vector, plus, each round, the mean feature profile of the revealed hits vs. non-hits, so the LLM can induce a decision rule and apply it.
- **A2** — real gene names (the C-arm).

Two design points matter. A1 is never shown the ML confidence score: were it shown, “reasoning” would collapse into copying the GBM. And **A0.5 exists because A1 is handed an ML-ranked pool that A0 is not** — without that control, A1 – A0 would confound “the features” with “being given a shortlist”, and on the two smallest screens (K562-Essential: 623 genes; Scharenberg22: 1029) that shortlist is itself a strong prior.

Table 6: Reasoning vs. recall. Hit ratio at round 5, mean over **5 seeds**, same backend, $T = 0$. The three right-hand columns isolate one ingredient each, holding the others fixed, and are *seed-paired* (mean of per-seed differences) ± 1 s.d. — not differences of means. A2 reproduces its committed 5-seed values on all nine screens (max drift 0.025).

Screen	hit@R5 (mean)					value of ... (seed-paired $\pm$ s.d.)		
	A0	A0.5	A1	A2	GBM	pool	features	name
						A0.5-A0	A1-A0.5	A2-A1
IFN- γ	0.180	0.160	0.175	0.190	0.183	-0.020 $\pm$.003	+0.015 $\pm$.008	+0.015 $\pm$.015
IL-2	0.351	0.316	0.343	0.347	0.314	-0.035 $\pm$.000	+0.027 $\pm$.007	+0.004 $\pm$.004
Sanchez21	0.088	0.090	0.090	0.095	0.087	+0.002 $\pm$.002	+0.001 $\pm$.005	+0.004 $\pm$.005
Sanchez21 (dn)	0.091	0.093	0.099	0.095	0.101	+0.002 $\pm$.002	+0.006 $\pm$.001	-0.003 $\pm$.005
Carnevale22	0.064	0.051	0.058	0.059	0.047	-0.013 $\pm$.001	+0.007 $\pm$.002	+0.001 $\pm$.002
Scharenberg22	0.469	0.531	0.584	0.449	0.449	+0.061 $\pm$.000	+0.053 $\pm$.011	- 0.135 $\pm$.011
Steinhart	0.075	0.048	0.065	0.145	0.090	-0.026 $\pm$.003	+0.017 $\pm$.008	+ 0.080 $\pm$.013
K562-Essential	0.286	0.318	0.597	0.549	0.476	+0.032 $\pm$.043	+ 0.279 $\pm$.014	-0.048 $\pm$.019
K562-GWPS	0.343	0.300	0.372	0.347	0.273	-0.043 $\pm$.000	+0.072 $\pm$.006	-0.024 $\pm$.006
Average	0.216	0.212	0.265	0.253	0.224	-0.005	+ 0.053	-0.012

Three results, in order of how much they surprised us:

1. **The shortlist is worth nothing** (-0.005 , negative on 5 of 9). A1’s gain is not an artefact of being handed ML-curated candidates.
2. **The features are worth $+0.053$, and positive on all nine screens** ($+0.001$ to $+0.279$), every per-screen effect exceeding its seed s.d. — the most consistent effect anywhere in this paper.
3. **The gene name is worth -0.012** . With every gene renamed to a meaningless identifier and the structural features supplied, the system *matches or slightly exceeds* the real-name arm (0.265 vs. 0.253). On average, *the LLM does not need to know that a gene is IRF4; it needs the structure*.

That average, however, hides a bimodality that is the whole point:

- **Steinhart: the name is worth $+0.080 \pm 0.013$** , the largest anywhere, and the features recover almost nothing ($+0.017 \pm 0.008$; A1 at 0.065 is even *below* A0 at 0.075, and the two error bars do not come close to overlapping). This is exactly the screen of Table 7: the one gain-of-function (CRISPRa) screen, where every DepMap feature is knockout-derived and does not transfer. The features we are able to hand the LLM are precisely the ones that are uninformative there, so nothing we supply substitutes for knowing what IRF4 does.
- **Scharenberg22: the name is worth -0.135 ± 0.011** — it actively hurts, consistent with the shuffle *improving* that screen.

So “reason, don’t recall” is achievable and essentially free on eight of nine screens, and impossible on the one screen where biology is doing real work. We read this as a limit of what our feature table *contains*, not as proof that reasoning cannot work: our features are anchor-relative PPI / co-expression / pathway overlap plus constraint and essentiality, and none of them encodes the specific T-cell biology Steinhart rewards. The obvious reading of the bullet above — that Steinhart is special *because* it is the gain-of-function screen — is tempting, and wrong: we tested it directly (next), and perturbation modality does not explain it.

One further calibration: **A0 (0.216) scores below the plain GBM (0.224)**. An LLM branch with nothing to say does not merely fail to help — at a 0.4 fusion weight it drags the hybrid below the ML model it was meant to improve.

Testing the explanation: is it the gain-of-function modality? (No.)

The story above rests on a sample of one: Steinhart is the only gain-of-function (CRISPRa) screen in the benchmark, so “gain-of-function makes the gene name matter” is a hypothesis fitted to a single point. We tested it. The same study that supplied our IL-2 and IFN- γ screens (Schmidt et al., 2022) ran *both* a CRISPRi and a CRISPRa genome-wide screen for each cytokine, so the CRISPRa arm is the matched gain-of-function counterpart of two screens we already have. We onboarded `IL2_crispra` and `IFNG_crispra` from that arm, keeping everything fixed except the modality: the same anchor-relative features (they are modality-agnostic — activating vs. silencing IL-2 biology shares the same anchors), the same hit convention, an activation-framed task prompt (“genes that, when *over-expressed via CRISPRa*, regulate IL-2 production”), and — because these large screens would otherwise route to the LLM-suppressed `ml_heavy` regime — the same fusion weight as Steinhart ($w_{\text{LLM}} = 0.4$).

If gain-of-function were the cause, the gene name should be worth $\approx +0.08$ here too. It is not:

Screen (gain-of-function, CRISPRa)	name value (A2–A1)	cf. its CRISPRi twin
<code>IL2_crispra</code>	$+0.002 \pm 0.003$	-0.050
<code>IFNG_crispra</code>	-0.004 ± 0.001	$+0.013$
<i>Steinhart</i> (reference)	$+0.080 \pm 0.013$	—

On both new gain-of-function screens the gene name is worth nothing, an order of magnitude below Steinhart. **Perturbation modality does not explain Steinhart’s name-dependence**. The corrected-prompt run barely moved these numbers from a first run with a placeholder prompt, so this is not an artefact of under-specifying the task; the name simply does not help here.

What *does* distinguish Steinhart is the shape of its hit set: 145 concentrated hits (0.8%) dominated by textbook exhaustion regulators (IRF4, TOX, ...) that a language model knows cold, against ~ 900 diffuse top-movers ($\sim 5\%$) on each CRISPRa cytokine screen, where recognizable names are a small minority. Our loss-of-function-trained prior also transfers poorly to CRISPRa hits (hit@R5 ≈ 0.12 vs. 0.35 on the CRISPRi twin), so the whole system is weak there. We leave the hit-set-concentration hypothesis untested. The operational consequence is clear enough, though: **we do not route by modality** — the structure-only reasoning mode stays an opt-in choice, not a default triggered by whether a screen is CRISPRa.

An accidental confirmation: the feature router already knew

The C-arm selects its feature set per screen. That router was tuned *empirically*, by trying feature sets and keeping whatever scored best, long before we understood why. It converged on excluding **all four DepMap features** for exactly two screens — Steinhart and K562-Essential — because including them *hurt*. The code comment recording this guessed at the reason (“DepMap disrupts online learning for GD2 biology”), and that guess was wrong: Steinhart does not measure GD2 biology at all (Table 1).

The real reason is the one above. **Every DepMap feature is derived from knockout screens**, and Steinhart is the one *gain-of-function* (CRISPRa) screen: it asks which genes make CAR-T cells resist exhaustion *when over-expressed*. A gene that is dispensable when deleted can

be potent when over-expressed — IRF4, among Steinhart’s strongest hits, is essential in only 6% of DepMap cell lines — so a knockout-derived prior is not merely uninformative there, it is actively misleading, and the router learned to discard it.

Reading the SHAP attributions back out confirms that this is what the model actually does:

Table 7: The ML model’s most-used features, per screen (mean |SHAP| over the genes it selected). DepMap essentiality dominates wherever it is available — and is absent exactly where a knockout-derived prior does not transfer.

Screen	Perturbation	Top features used by the ML model
IFN- γ , IL-2, Carnevale22, Scharenberg22, K562-GWPS	loss-of-function	depmap_frac_ess (rank 1), then PPI / pLI
Steinhart	gain-of-function (CRISPRa)	PPI, STRING degree, pathway counts — no DepMap feature at all
K562-Essential	loss-of-function	PPI, STRING degree — no DepMap feature

This closes the loop on why Steinhart is the screen where the LLM matters most: it is precisely the screen where the ML model’s single strongest feature family has been removed, leaving it with weak topological signal — and the LLM’s T-cell biology is what fills the gap. (K562-Essential drops DepMap for a different reason: it *is* a curated essentiality screen, so a pan-cancer essentiality prior is close to a restatement of the label and adds nothing.)

Does giving the LLM tools help? (No.)

The natural next step from a pipeline is a *tool-using agent* in the BioDiscoveryAgent / PerTurboAgent mould: let the LLM plan, call an ML-ranking tool and a pathway-enrichment tool, and commit a batch itself. We implemented exactly that (a task-specialised harness with a JSON action protocol; `ml_rank` / `enrich` / `finish`; no agent framework) and benchmarked it against the pipeline on identical screens and budgets.

Table 8: Tool-using agent vs. the deterministic pipeline (hit ratio @ R5). The agent plans its own rounds and calls tools; the pipeline fuses online ML with LLM reasoning under a fixed policy. The agent is *worse on 8 of 9 screens*.

Dataset	Tool-using agent	Pipeline (C)	Δ
Scharenberg22	0.582	0.465	+0.116
Carnevale22	0.044	0.056	−0.012
IFN- γ	0.163	0.190	−0.027
Sanchez21	0.059	0.087	−0.029
Sanchez21 (dn)	0.049	0.099	−0.050
Steinhart	0.100	0.159	−0.059
IL-2	0.264	0.347	−0.083
K562-GWPS	0.209	0.347	−0.138
K562-Essential	0.413	0.556	−0.143
Average	0.209	0.256	−0.047




Figure 4: Per-screen delta of the tool-using agent against the pipeline. It wins on exactly one screen and loses worst precisely where the ML signal is strongest (K562-Essential, K562-GWPS) — given freedom, the LLM *overrides* a well-calibrated model.

The agent wins on exactly one screen (Scharenberg22, +0.116), where its runtime pathway-enrichment tool pays off, and loses most on the two genome-wide screens where the ML prior is strongest. The failure mode is instructive: given the freedom to act, the LLM overrides a model that is better calibrated than it is. **Constraining the LLM to a fixed fusion policy is not a limitation of the pipeline — it is the reason the pipeline wins.**

Transplanting the agent’s one good idea

Since enrichment was the agent’s only advantage, we transplanted it into the pipeline (Enrichr over the hits revealed so far, injected into the LLM prompt each round, gated to significant and coherent pathways), and separately tested an externalised *skill library* distilled from past experiments.

Table 9: Agent-style augmentations of the pipeline, as **paired** deltas over the same runs. Both are within noise of zero.

Augmentation	C	C + augmentation	Δ (paired)
Runtime pathway enrichment (gated)	0.260	0.263	+0.003
Externalised skill library	0.258	0.259	+0.001

Runtime enrichment captures the Scharenberg22 gain (+0.061 on that screen) without the collateral damage the free agent suffered, but nets only +0.003 overall: on genome-wide essentiality screens the enriched terms are highly significant yet *non-selective* (ribosome, proteasome, spliceosome), so “prioritise genes in these pathways” barely narrows a 9,000-gene pool. The skill library adds +0.001. Together with the memory ablation (+0.000), this is a consistent negative result about *externalisation*: memory, skills and retrieved pathways all add ≈ 0 once the model already has its parametric knowledge plus the hits it has observed.

Explainability: what is the LLM actually contributing?

The ablations say *how much* each component is worth; they cannot say *where* the value comes from. We therefore instrumented the arm to record, for every selected gene, a counterfactual attribution against the ML model’s own top- k :

- **both** — the LLM named it *and* the ML would have selected it anyway;
- **LLM-only** — the LLM named it and the ML would *not* have selected it (this is the LLM’s marginal contribution);
- **ML-only** — the LLM did not name it; it entered on ML score alone.

Only genes the LLM actually named count as LLM picks. The LLM branch pads its batch from the static ranker (Table 3[†]), and crediting that padding to the LLM would make “ML and LLM agree” partly mean “two ML models agree”. An earlier version of this analysis made exactly that error: it reported an agreement bucket of $n = 800$ at a 21.1% hit rate, of which $\approx 87\%$ was static-ranker padding. Excluding it sharpens the result considerably.

Pairing each pick with its revealed outcome gives the hit rate *per source* (9 screens $\times$ 5 rounds, $\approx 4,200$ selected genes). K562-GWPS is reported separately: it is the one screen routed through the two-stage policy (the ML shortlists 384 candidates and the *LLM* makes the final selection), so the “ML-only” category cannot exist there and pooling it would corrupt the rates.

Table 10: Hit rate by attribution source, over the eight weighted-fusion screens. *Pooled* counts every selected gene once; *macro* averages the per-screen rates (screens with ≥ 10 picks in that category). Genuine agreement is rare — only 101 of $\approx 4,200$ picks — but it is by far the strongest signal, at roughly **three times** the ML’s own hit rate. The LLM’s unilateral picks are *worse* than the ML’s when pooled and about equal when macro-averaged, so the defensible claim is that they are *no better*.

Source	Genes selected	Hits	Hit rate (pooled)	Hit rate (macro)
ML + LLM agree	101	47	46.5% $\pm$ 9.7	41.5%
ML only	3,307	487	14.7% $\pm$ 1.2	14.9%
LLM only	752	61	8.1% $\pm$ 2.0	13.1%




Figure 5: Left: what each round’s batch is made of, pooled over the eight weighted-fusion screens. Right: the hit rate of each source. Genes that *both* components endorse hit at roughly twice the rate of genes the LLM introduces on its own.

The LLM is a verifier, not a proposer. The LLM is a *poor proposer*: left to itself it picks genes that hit at 8.1%, well below the ML’s own 14.7%. But it is a *strong verifier*: when it independently names a gene the ML also ranks in its top- k , that gene hits at 46.5% — roughly three times the ML’s rate. Such agreement is rare (101 of $\approx 4,200$ picks, i.e. 2.4%), which is precisely why the component is worth so little on average and so much where it fires. This single measurement explains the two results above that otherwise look unrelated:

- why **removing the LLM costs only 0.008** — the genes it proposes on its own are *worse* than what the ML would have picked instead, so losing them costs nothing; what is lost is a rare but highly informative agreement signal;
- why **giving the LLM tools makes it worse** (-0.047 , Table 8) — a free agent selects mostly on its own judgement, which is precisely the weakest of the three sources.

The design implication is the opposite of the prevailing “more agency” direction: the LLM should be used to *re-weight* a calibrated model’s candidates, not to choose freely.




Figure 6: Left: features driving the ML score over the genes actually selected (LightGBM SHAP contributions; DepMap essentiality and PPI connectivity dominate). Right: calibration — the top-scoring quintile of tested genes hits about twice as often as the bottom quintile, though the curve is flat in the middle and the range is narrow because only *selected* genes are ever tested.

Limitations

We cannot separate the LLM’s biological knowledge from memorisation. This is the most important caveat on every LLM result above. Seven of the nine screens are published, so their hit genes may sit in the model’s training data, and a gene the model “knows” is a hit is indistinguishable — to our measurements — from a gene it reasons its way to.

It is tempting to appeal to the Steinhart screen as a clean control: BioDiscoveryAgent introduced it precisely because it was *unpublished*, stating that it is therefore “not part of the language model’s training data”. That argument was sound *for their models* (Claude 3.5 Sonnet, whose training cutoff precedes their release). **It does not transfer to ours.** BDA released the screen in their public repository (`ground_truth_Steinhart_crispra_GD2_D22.csv` and the corresponding `topmovers` file) in 2024, whereas our LLM is Claude Haiku 4.5, released in late 2025. We therefore cannot claim the screen is outside our model’s training data, and we withdraw the otherwise attractive inference that “the LLM contributes most on the one screen it provably cannot have memorised”.

We judge the practical leakage risk to be low — the hit set exists only as an 18,000-row table of effect sizes and a binary `.npy` array, not the kind of artefact language models reliably memorise — but low is not zero, and the strong claim is not available to us. Note also that the gene-name-shuffle ablation (-0.084 on Steinhart) does *not* settle this: shuffling gene names destroys memorised gene-hit associations and genuine biological knowledge equally, so it cannot distinguish them. Settling the question requires a screen released *after* the model’s training cutoff; none of our nine qualifies.

Other limitations.

- **Not a reproduction of BioDiscoveryAgent.** We use their data and their non-essential convention, but never ran their agent; our numbers are not comparable to theirs (we add a cross-experiment ML prior they do not have; Table 1 and the following section).
- **The B baseline is padded**, on some screens heavily (86% on K562-Essential); it is not a pure-LLM baseline (Table 3[†]).
- **Genuine ML/LLM agreement is rare** ($n = 101$), so its hit rate carries a wide interval ($46.5\% \pm 9.7$). The *ordering* of the three sources is robust; the exact rate is not.

- **The attribution excludes K562-GWPS**, whose two-stage route lets the LLM make the final selection, so an “ML-only” pick cannot exist there by construction.
- **Calibration is measured only over genes the agent chose to test**, so the score range is narrow by construction and the curve is flat in its interior.
- **Nine screens, five seeds, no prospective wet-lab validation.** The arm is stochastic, which is why every component claim is reported as a *paired* delta.